Scalable Variational Inference

Scalable Variational Inference

People used to thing that Bayes is for small datasets

Scalable Variational Inference

People used to thing that Bayes is for small datasets

- Too slow for big data

Scalable Variational Inference

People used to thing that Bayes is for small datasets

- Too slow for big data
- Not very beneficial anyway

Scalable Variational Inference

People used to thing that Bayes is for small datasets

- Too slow for big data
- Not very beneficial anyway

Things changed when Bayes met Deep Learning

Scalable Variational Inference

- Understand why and how combine Deep Learning and Bayesian methods

Scalable Variational Inference

- Understand why and how combine Deep Learning and Bayesian methods

- Learn how to synthesize images with VAE

Scalable Variational Inference

- Understand why and how combine Deep Learning and Bayesian methods

- Learn how to synthesize images with VAE

- Learn state-of-the-art Bayesian neural networks and their applications

Unbiased estimates

Unbiased estimates

$$\mathbb{E}_{p(x)}\, f(x)$$

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s)$$

$$x_s \sim p(x)$$

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s)$$

$$x_s \sim p(x)$$

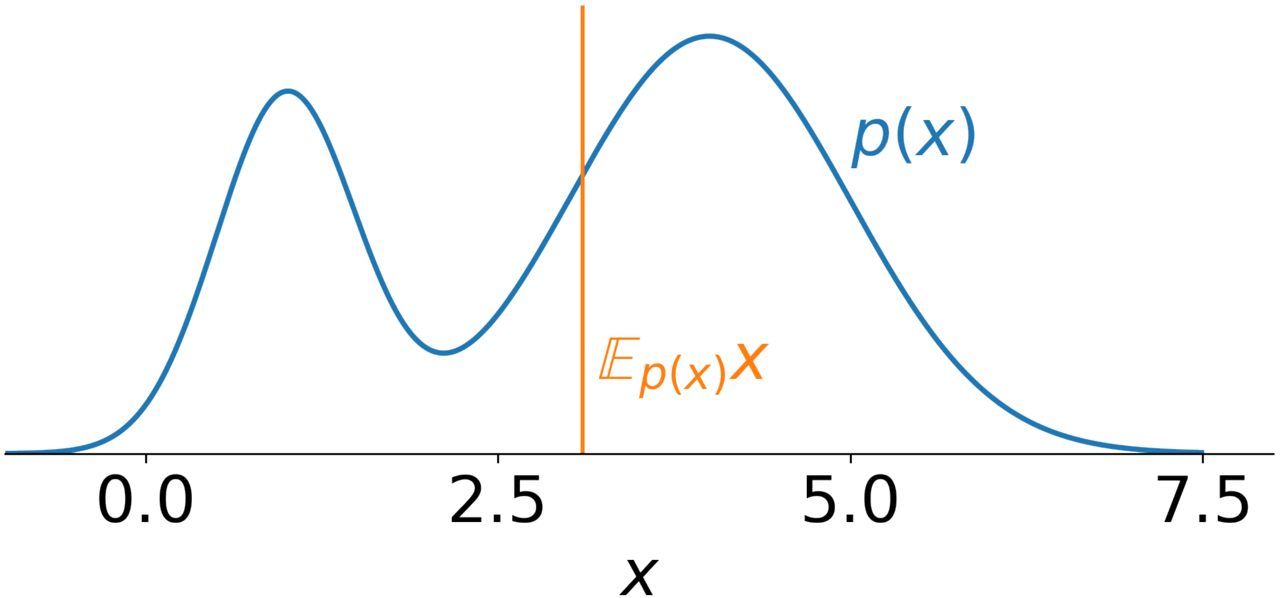

Unbiased estimates

$$\mathbb{E}_{p(x)}\, f(x) \approx \frac{1}{M}\sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

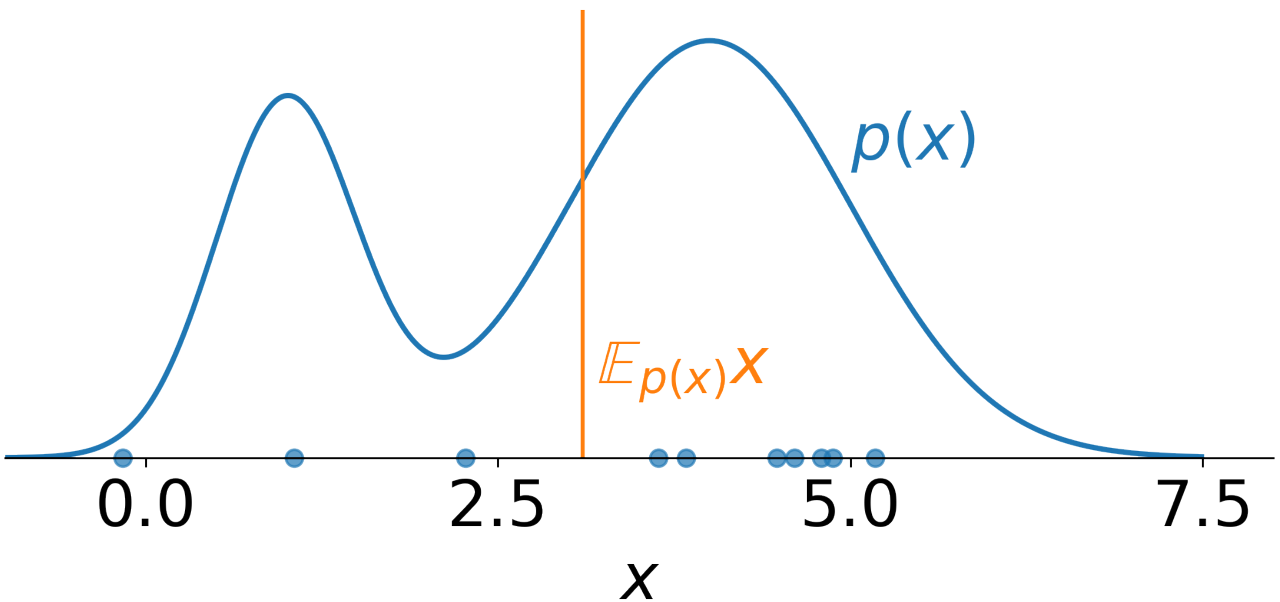

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

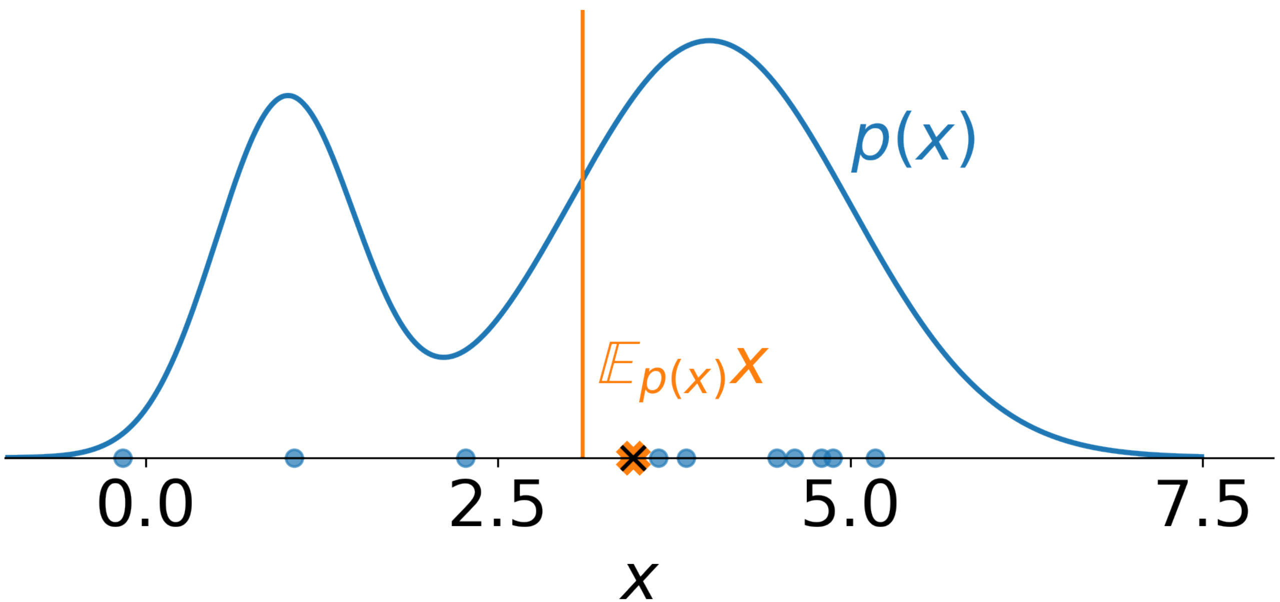

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

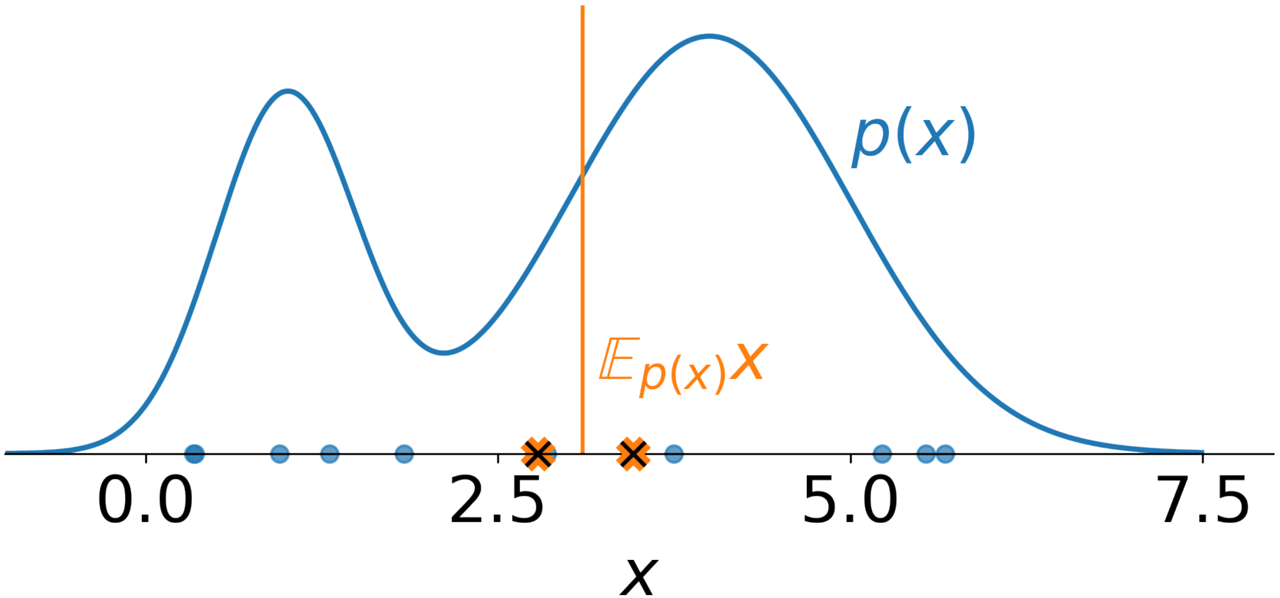

Unbiased estimates

$$\mathbb{E}_{p(x)}\, f(x) \approx \frac{1}{M}\sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

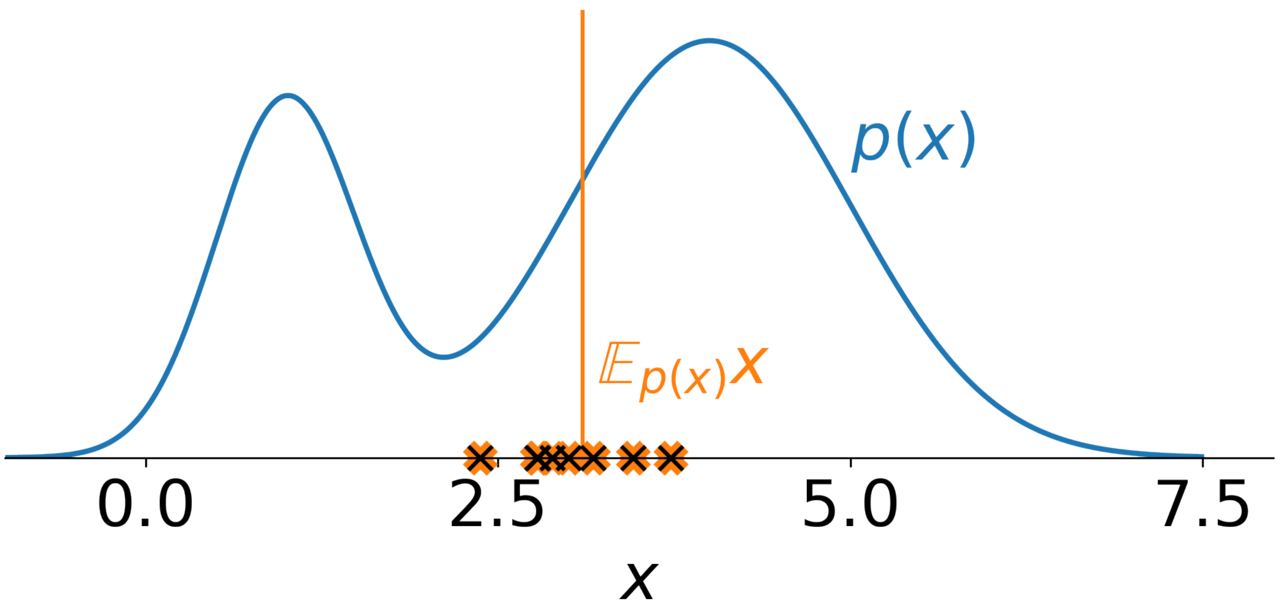

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

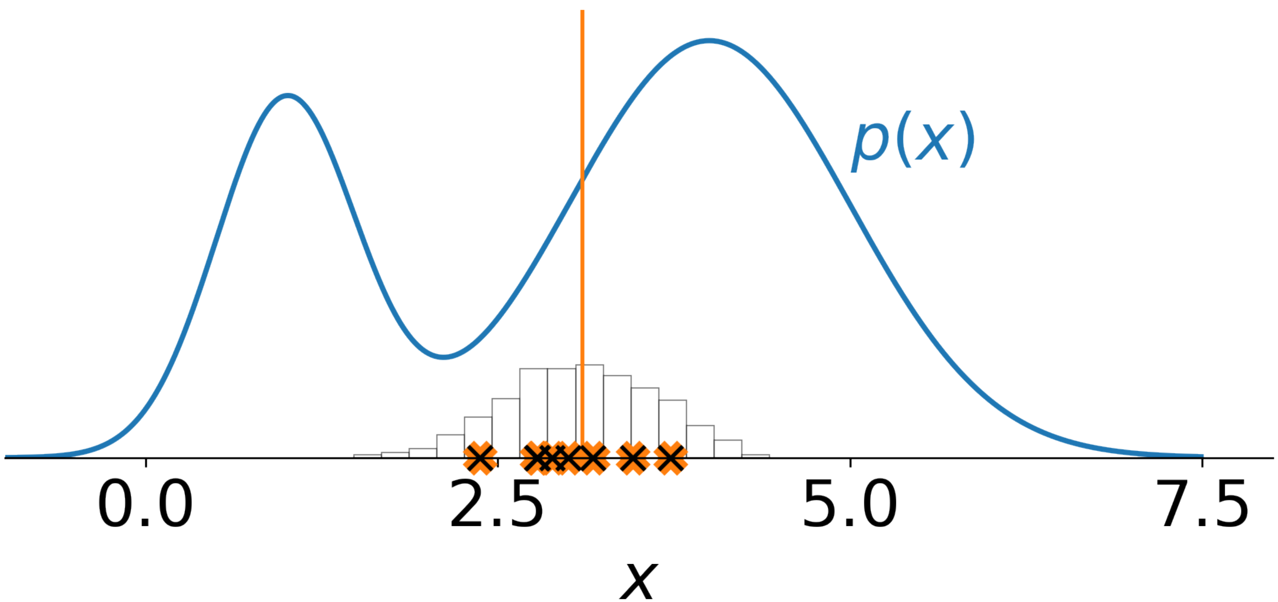

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

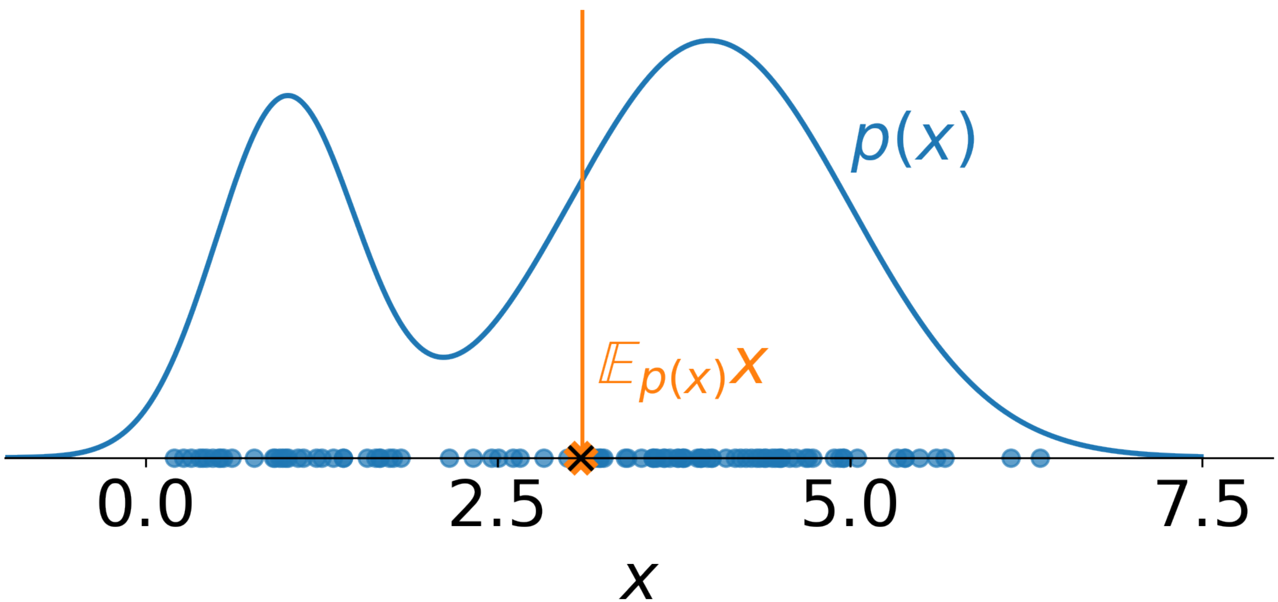

Unbiased estimates

$$\mathbb{E}_{p(x)} f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

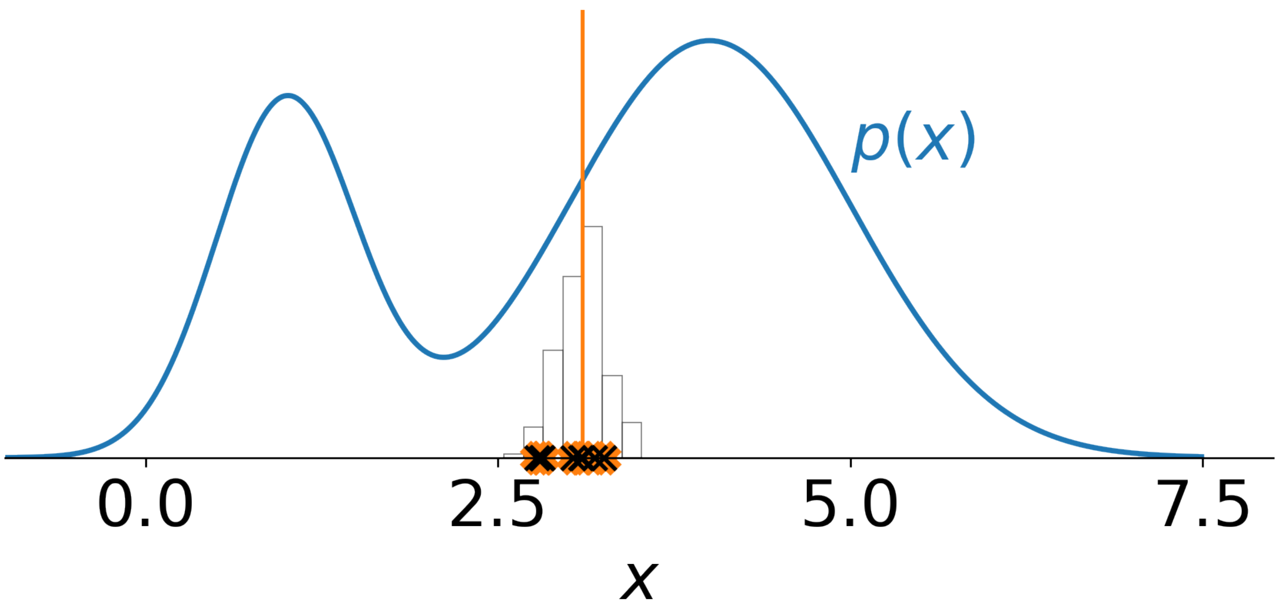

Unbiased estimates

$$\mathbb{E}_{p(x)}\, f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

$$x_s \sim p(x)$$

$$\mathbb{E}_{p(x)}\, R = \mathbb{E}_{p(x)}\, f(x)$$

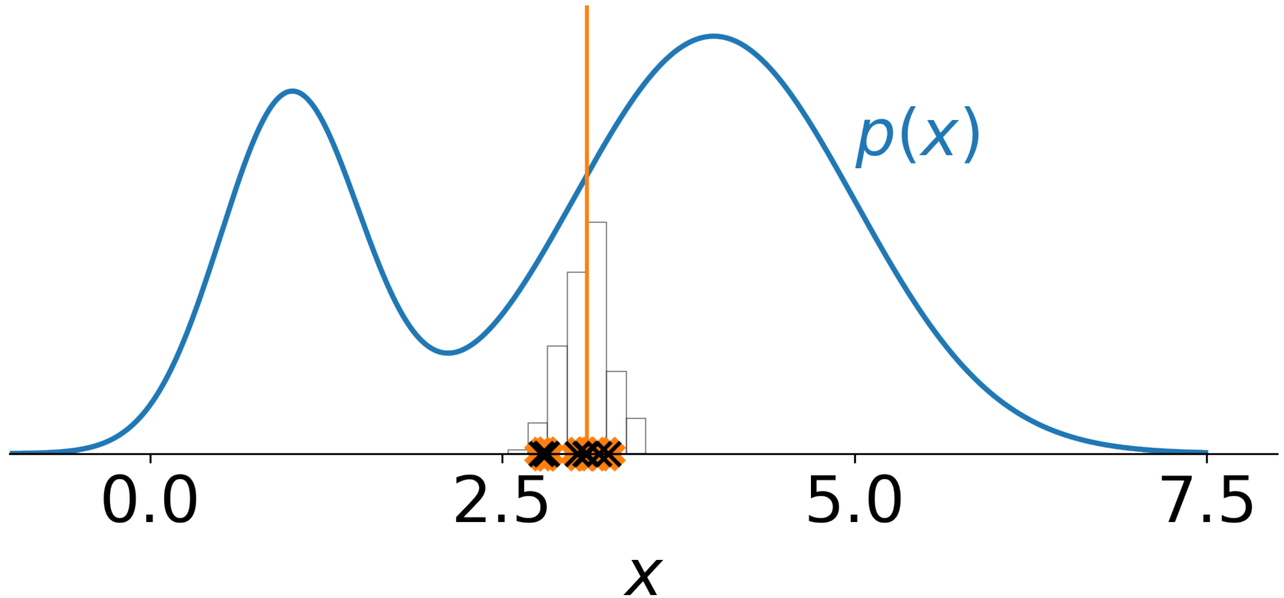

Unbiased estimates

$$\log\left(\mathbb{E}_{p(x)}\, f(x)\right) \overset{?}{\approx} \log\left(\frac{1}{M}\sum_{s=1}^{M} f(x_s)\right) = G$$

$$x_s \sim p(x)$$

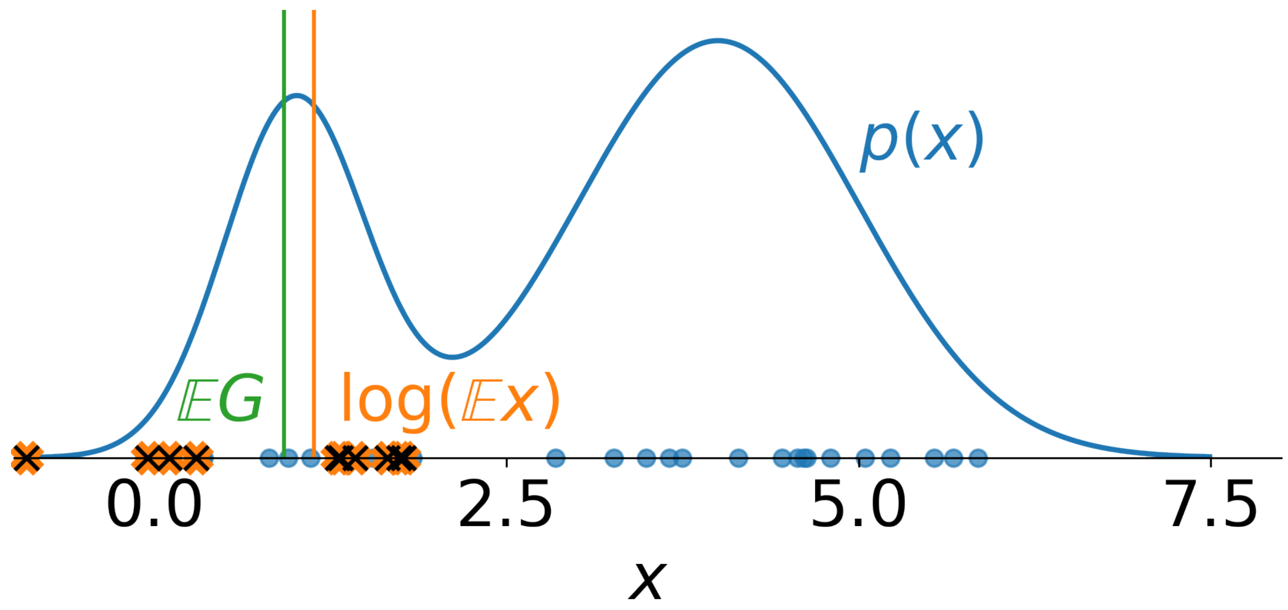

Unbiased estimates

$$\log\left(\mathbb{E}_{p(x)} f(x)\right) \overset{?}{\approx} \log\left(\frac{1}{M}\sum_{s=1}^{M} f(x_s)\right) = G$$

$$x_s \sim p(x)$$

$$\mathbb{E}_{p(x)} G \neq \log\left(\mathbb{E}_{p(x)} f(x)\right)$$

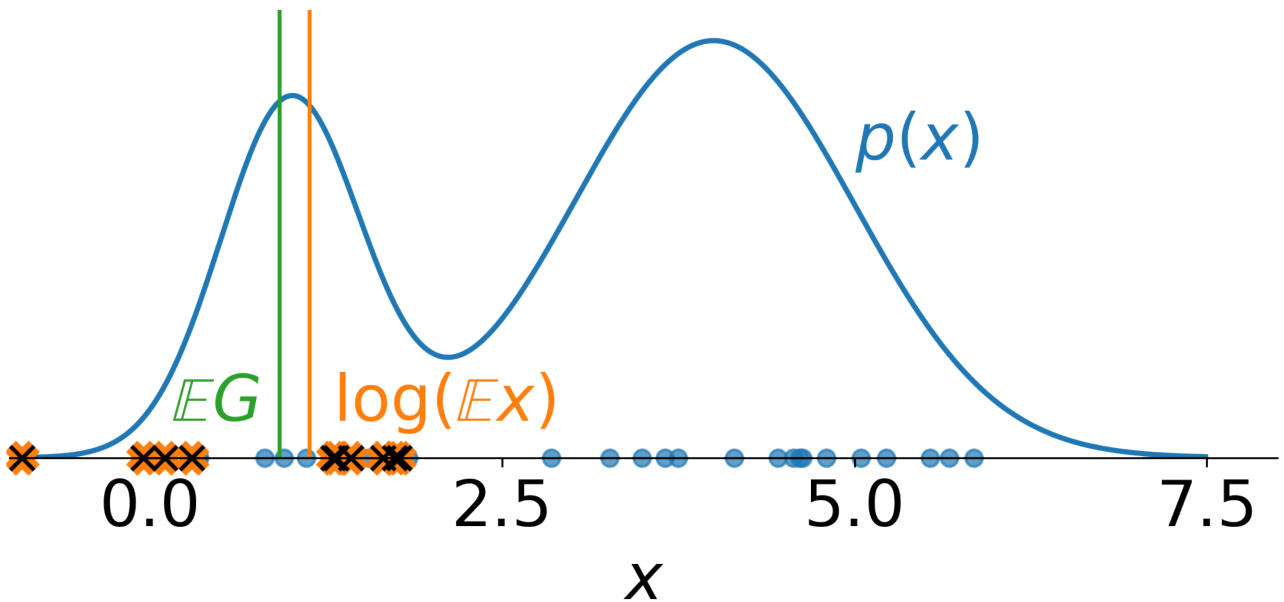

Unbiased estimates

- Estimator called unbiased if its expected value equals to thing it estimates

Unbiased estimates

- Estimator called unbiased if its expected value equals to thing it estimates

- This, is unbiased estimator:

$$\mathbb{E}_{p(x)}\, f(x) \approx \frac{1}{M} \sum_{s=1}^{M} f(x_s) = R$$

others may look unbiased, but you have to check